deeplearning.ai

Comparing to human-level performance

Improving your model performance

The two fundamental assumptions of supervised learning

1. You can fit the training set pretty well.

2. The training set performance generalizes pretty well to the dev/test set.

Reducing (avoidable) bias and variance

Human-level

Train bigger model

Train longer/better optimization algorithms

Training error

NN architecture/hyperparameters search

More data

Dev error

Regularization

NN architecture/hyperparameters search

Andrew Ng